


Neural Networks - Part II

Arash Marioriyad


19 Khordad 1405

CE 417-Artificial Intelligence
Sharif Univ. of Tech.

Most Slides have been adopted from Berkeley CS188 & SBU CI
course Dr. Malek & SUT CE417 and CE717





Perceptron





AND & OR Networks



-



Two-Class Perceptron



What Binary Threshold Neurons Cannot Do



- Cannot compute an XOR



Learning: Binary Perceptron

- ▶ Given N training instances $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})$
 - ▶ $y^{(n)} = +1$ or -1

- Initialize $\mathbf{w}$
- Cycle through the training instance
- While more classification errors

- For $i = 1 \dots N_{train}$
 - $\hat{y}^{(i)} = \text{sign}(\mathbf{w}^T \mathbf{x}^{(i)})$
 - If $\hat{y}^{(i)} \neq y^{(i)}$
 - $\mathbf{w} = \mathbf{w} + y^{(i)} \mathbf{x}^{(i)}$

If instance misclassified:

– If instance is positive class

$$\mathbf{w} = \mathbf{w} + \mathbf{x}$$

– If instance is negative class


$$\mathbf{w} = \mathbf{w} - \mathbf{x}$$




How to get probabilistic decisions?

- Perceptron scoring: $z = w \cdot f(x)$
- If $z = w \cdot f(x)$ very **positive** $\Rightarrow$ want probability going to 1
- If $z = w \cdot f(x)$ very **negative** $\Rightarrow$ want probability going to 0
- Sigmoid function (logistic function)


$$\phi(z) = \frac{1}{1 + e^{-z}}$$



- We can use sigmoid as an activation Function (instead of step function)



Two-Class Logistic Regression



Activation
Function




Logistic regression: loss function

Assume we have two classes with labels 0 and 1


$$\hat{\mathbf{w}} = \underset{\mathbf{w}}{\operatorname{argmin}} J(\mathbf{w})$$

$$\begin{aligned} J(\mathbf{w}) &= \\ &= \sum_{i=1}^n -y^{(i)} \log \left(\sigma(\mathbf{w}^T \mathbf{x}^{(i)}) \right) - (1 - y^{(i)}) \log \left(1 - \sigma(\mathbf{w}^T \mathbf{x}^{(i)}) \right) \end{aligned}$$


$J(\mathbf{w})$ is convex w.r.t. parameters.



Logistic Regression for 3-way classification




$$\begin{aligned} z_i &= w_i \cdot f \\ &= \sum_j w_{ji} \cdot f_j \end{aligned}$$




Logistic Regression for 3-way classification

- Deep Neural Network = Also Learn the Features End-to-End





Deep Neural Network for 3-way classification




...

Layer L




Deep Neural Network for 3-way classification




$$h_1^{(1)} = \phi \left(w_1^{(1)} \cdot x \right) = \phi \left(\sum_j w_{ji}^{(1)} \cdot x_j \right)$$

ϕ = activation function



Deep Neural Network for 3-way classification




$$h_1^{(2)} = \phi \left(w_1^{(2)} \cdot h^{(1)} \right) = \phi \left(\sum_j w_{ji}^{(2)} \cdot h_j^{(1)} \right)$$


ϕ = activation function



Deep Neural Network for 3-way classification




...




$$h_i^{(l)} = \phi\left(\sum_j w_{ji}^{(l)} \cdot h_j^{(l-1)}\right)$$

ϕ = activation function


- Neural network with L layers
- $h^{(l)}$: activations at layer l
- $w^{(l)}$: weights taking activations from layer l-1 to layer l



Deep Neural Network for 3-way classification




...




$$h^{(l)} = \phi(h^{(l-1)} \times W^{(l)})$$


ϕ = activation function



Deep Neural Network for 3-way classification




- Sometimes also called *Multi-Layer Perceptron (MLP)* or *Feed-Forward Network (FFN)*



Non Linearity is Important

- Common Activation Functions


Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$




MLP Universal Approximator

- ▶ A feed-forward network with a single hidden layer and linear outputs can approximate any continuous function on a compact domain to an arbitrary accuracy
 - ▶ under mild assumptions on the activation function
 - ▶ e.g., sigmoid activation functions (Cybenko, 1989)
 - ▶ when sufficiently large (but finite) number of hidden units is used


$$F_k(x) = \sum_{j=1}^M w_{kj}^{[2]} \phi \left(\sum_{i=0}^d w_{ji}^{[1]} x_i \right)$$

- ▶ It is of greater theoretical interest than practical
 - ▶ the construction of such a network requires the nonlinear activation functions and weight values which are unknown



XOR Example

$$f = \text{XOR}(x_1, x_2) = \text{OR}(\text{AND}(x_1, \bar{x}_2), \text{AND}(\bar{x}_1, x_2))$$



General Boolean Functions

Truth table shows all input combinations
for which output is 1

Truth Table

X ₁	X ₂	X ₃	X ₄	X ₅	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + \\ X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + X_1\bar{X}_2X_3X_4X_5 + X_1X_2\bar{X}_3\bar{X}_4X_5$$





General Boolean Functions (cont.)

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

Truth table shows all input combinations
for which output is 1

$$y = \bar{X}_1 \bar{X}_2 X_3 X_4 \bar{X}_5 + \bar{X}_1 X_2 \bar{X}_3 X_4 X_5 + \bar{X}_1 X_2 X_3 \bar{X}_4 \bar{X}_5 + \\ X_1 X_2 X_3 X_4 X_5 + X_1 \bar{X}_2 X_3 X_4 X_5 + X_1 X_2 \bar{X}_3 \bar{X}_4 X_5$$





General Boolean Functions (cont.)

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

Truth table shows all input combinations
for which output is 1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + X_1X_2X_3X_4X_5 + X_1X_2\bar{X}_3\bar{X}_4X_5$$





General Boolean Functions (cont.)

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

Truth table shows all input combinations
for which output is 1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + X_1\bar{X}_2X_3X_4X_5 + X_1X_2X_3X_4X_5$$




General Boolean Functions (cont.)

Truth table shows all input combinations
for which output is 1

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + \\ \boxed{X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5} + X_1\bar{X}_2X_3X_4X_5 + X_1X_2\bar{X}_3\bar{X}_4X_5$$





General Boolean Functions (cont.)

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

Truth table shows all input combinations
for which output is 1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + \\ X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + \boxed{X_1\bar{X}_2X_3X_4X_5} + X_1X_2\bar{X}_3\bar{X}_4X_5$$





General Boolean Functions (cont.)

Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

Truth table shows all input combinations
for which output is 1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + \\ X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + X_1\bar{X}_2X_3X_4X_5 + X_1X_2\bar{X}_3\bar{X}_4X_5$$





General Boolean Functions (cont.)

Truth table shows all input combinations
for which output is 1


Truth Table

X_1	X_2	X_3	X_4	X_5	Y
0	0	1	1	0	1
0	1	0	1	1	1
0	1	1	0	0	1
1	0	0	0	1	1
1	0	1	1	1	1
1	1	0	0	1	1

$$y = \bar{X}_1\bar{X}_2X_3X_4\bar{X}_5 + \bar{X}_1X_2\bar{X}_3X_4X_5 + \bar{X}_1X_2X_3\bar{X}_4\bar{X}_5 + X_1\bar{X}_2\bar{X}_3\bar{X}_4X_5 + X_1\bar{X}_2X_3X_4X_5 + X_1X_2\bar{X}_3\bar{X}_4X_5$$



General Functions





Can now be composed into “networks” to compute arbitrary classification “boundaries”



- Build a network of units with a single output that fires if the input is in the coloured area





General Functions (cont.)





General Functions (cont.)





General Functions (cont.)





General Functions (cont.)





General Functions (cont.)





General Functions (cont.)



More Complex Decision Boundaries



More Complex Decision Boundaries (cont.)



MLP with Different Number of Layers



MLP with unit step activation function

Decision region found by an output unit.

Structure	Type of Decision Regions	Interpretation	Example of region
Single Layer (no hidden layer)	Half space	Region found by a hyper-plane	
Two Layer (one hidden layer)	Polyhedral (open or closed) region	Intersection of half spaces	
Three Layer (two hidden layers)	Arbitrary regions	Union of polyhedrals	




Exercise: Compose This with One Hidden Layer



Expressive Power of MLPs

- ▶ MLPs are universal Boolean function
- ▶ MLPs are universal classifiers
- ▶ MLPs are universal function approximators
- ▶ An MLP with two (or even one) hidden layers can approximate anything to arbitrary precision
 - ▶ But could be exponentially or even infinitely wide in its inputs size




Learning Problem in Summary

- Given: the architecture of the network
- Training data: A set of input-output pairs

$$(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), (\mathbf{x}^{(2)}, \mathbf{y}^{(2)}), \dots, (\mathbf{x}^{(N)}, \mathbf{y}^{(N)})$$

- We want to find the function g (model) on the input to get the output
- We consider a neural network as a parametric function $g(\mathbf{x}; \mathbf{W})$
- We need a **loss function** to show how penalizes the obtained output $g(\mathbf{x}; \mathbf{W})$ when the desired output is $\mathbf{y}$


$$\frac{1}{N} \sum_{n=1}^N \text{loss}(g(\mathbf{x}^{(n)}; \mathbf{W}), \mathbf{y}^{(n)})$$



Choosing Cost Function: Examples

- > Regression problem (multiple outputs)

SSE Loss



$$E = \sum_{n=1}^N E_n$$


$$E_n = \sum_{k=1}^K \left(\hat{y}_k^{(n)} - y_k^{(n)} \right)^2$$

- > Classification problem

Cross Entropy Loss

$$loss_n = \sum_{k=1}^K -y_k^{(n)} \log \hat{y}_k^{(n)}$$

Output is found by a softmax layer $\hat{y}_k = \frac{e^{z_k}}{\sum_{k=1}^K e^{z_k}}$




Deep Neural Network Training

- Do the maximum likelihood estimation is a good Idea ...
- Find a set of parameters which results in the best performance in training data
- In general, cannot always take derivative and set to 0
- Use numerical optimization!
- **GRADIENT DESCENT!**
- Hill climbing in continuous space



The Main Issue of Using Gradient Descent for NNs

- The loss function (L) itself is a function of model g
- The model g itself is a complex function of weights (parameters) $W^{(1)}$, $W^{(2)}$, ..., $W^{(L)}$
- We want the gradient of the loss function w.r.t to the parameters



Gradient in n Dimensions

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \dots \\ \frac{\partial L}{\partial w_n} \end{bmatrix}$$



Optimization Procedure: Gradient Descent

```
Init  $w$   
for iter = 1, 2, ...  
     $w \leftarrow w - \alpha \cdot \nabla L(w)$ 
```

L is the loss function

α : learning rate



Univariate Chain Rule

$$\frac{d}{dt}f(x(t)) = \frac{df}{dx} \frac{dx}{dt}.$$

- Example:

Computing the loss:

$$z = wx + b$$

$$y = \sigma(z)$$

$$\mathcal{L} = \frac{1}{2}(y - t)^2$$


Computing the derivatives:

$$\frac{d\mathcal{L}}{dy} = y - t$$

$$\frac{d\mathcal{L}}{dz} = \frac{d\mathcal{L}}{dy} \sigma'(z)$$


$$\frac{\partial \mathcal{L}}{\partial w} = \frac{d\mathcal{L}}{dz} x$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{d\mathcal{L}}{dz}$$



Multivariate Chain Rule

$$\frac{d}{dt} f(x(t), y(t)) = \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt}$$



Example:


$$f(x, y) = y + e^{xy}$$

$$x(t) = \cos t$$

$$y(t) = t^2$$



Plug in to Chain Rule:

$$\begin{aligned} \frac{df}{dt} &= \frac{\partial f}{\partial x} \frac{dx}{dt} + \frac{\partial f}{\partial y} \frac{dy}{dt} \\ &= (ye^{xy}) \cdot (-\sin t) + (1 + xe^{xy}) \cdot 2t \end{aligned}$$





Computation Graph

- We can diagram out the computations using a computation graph.
- The nodes represent all the inputs and computed quantities, and the edges represent which nodes are computed directly as a function of other nodes.



Backpropagation Algorithm

- Let $v_1, \dots, v_N$ be a topological ordering of the computation graph (i.e. parents come before children.)
- v_N denotes the variable we are trying to compute derivatives of (e.g. loss)



Backpropagation Example 1

Gradient of $g(w_1, w_2, w_3) = w_1^4 w_2 + 5w_3$ at $w_1 = 2, w_2 = 3, w_3 = 2$

Think of g as a composition of many functions

- Then, we can use the chain rule to compute the gradient

$$g = b + c$$

$$\frac{\partial g}{\partial b} = 1, \frac{\partial g}{\partial c} = 1$$

$$b = a \times w_2$$



$$\frac{\partial g}{\partial a} = \frac{\partial g}{\partial b} \frac{\partial b}{\partial a} = 1 \cdot w_2 = 3 \quad \frac{\partial g}{\partial w_2} = \frac{\partial g}{\partial b} \frac{\partial b}{\partial w_2} = 1 \cdot a = 16$$

$$a = w_1^4$$

$$\frac{\partial g}{\partial w_1} = \frac{\partial g}{\partial a} \frac{\partial a}{\partial w_1} = 3 \cdot 4w_1^3 = 96$$


$$c = 5w_3$$

$$\frac{\partial g}{\partial w_3} = \frac{\partial g}{\partial c} \frac{\partial c}{\partial w_3} = 1 \cdot 5 = 5$$



Backpropagation Example 2

- Build a computation graph and apply chain rule: $f(x) = g(h(x)) \quad f'(x) = h'(x) \cdot g'(h(x))$
- Example: neural network with quadratic loss: $L(a_2, y^*) = \frac{1}{2}(a_2 - y^*)^2$ and ReLU activations $g(z) = \max(0, z)$
- $a_2 = g_2(w_2 * g_1(w_1 * x))$



Main Goal of Backpropagation


- How to compute

$$\frac{d \text{loss}(\mathbf{o}, \mathbf{y})}{dw_{ji}^{[k]}}$$



Backpropagation: Notation

- ▶ $\mathbf{a}^{[0]} \leftarrow \text{Input}$
- ▶ $\text{output} \leftarrow \mathbf{a}^{[L]}$



Backpropagation: Last Layer Gradient

For squared error loss:



$$loss = \frac{1}{2} \sum_j (o_j - y_j)^2$$

$o_j = a_j^{[L]}$

$$\frac{\partial loss}{\partial a_j^{[L]}} = (a_j^{[L]} - y_j)$$

$$\frac{\partial loss}{\partial w_{ji}^{[L]}} = ?$$

$$a_i^{[L]} = f(z_i^{[L]})$$
$$z_j^{[L]} = \sum_{i=0}^M w_{ji}^{[L]} a_i^{[L-1]}$$



Backpropagation: Last layer gradient

For squared error loss:

$$\text{loss} = \frac{1}{2} \sum_j (\hat{y}_j - y_j)^2$$

$$\hat{y}_j = a_j^{[L]}$$

$$\frac{\partial \text{loss}}{\partial a_j^{[L]}} = (a_j^{[L]} - y_j)$$

$$\frac{\partial \text{loss}}{\partial w_{ji}^{[L]}} = \frac{\partial \text{loss}}{\partial a_j^{[L]}} \frac{\partial a_j^{[L]}}{\partial w_{ji}^{[L]}}$$

$$\frac{\partial a_j^{[L]}}{\partial w_{ji}^{[L]}} = f'(z_j^{[L]}) \frac{\partial z_j^{[L]}}{\partial w_{ji}^{[L]}} = f'(z_j^{[L]}) a_i^{[L-1]}$$

$$\frac{\partial \text{loss}}{\partial w_{ji}^{[L]}} = \frac{\partial \text{loss}}{\partial a_j^{[L]}} f'(z_j^{[L]}) a_i^{[L-1]}$$


Activation output of
neuron number j in
layer L

Pre activation value (before applying
the activation function) of neuron
number j in layer L

$$a_j^{[L]} = f(z_j^{[L]})$$

$$z_j^{[L]} = \sum_{i=0}^M w_{ji}^{[L]} a_i^{[L-1]}$$

**f is the
activation
function**



$$\frac{\partial \text{loss}}{\partial a_j^{[L]}}$$




$$\frac{\partial \text{loss}}{\partial w_{ji}^{[L]}}$$



Previous Layers Gradients

$$\frac{\partial \text{loss}}{\partial w_{ji}^{[l]}} = \frac{\partial \text{loss}}{\partial z_j^{[l]}} \frac{\partial z_j^{[l]}}{\partial w_{ji}^{[l]}} = \frac{\partial \text{loss}}{\partial z_j^{[l]}} a_i^{[l-1]}$$


$$\begin{aligned} \frac{\partial \text{loss}}{\partial z_i^{[l-1]}} &= \frac{\partial a_i^{[l-1]}}{\partial z_i^{[l-1]}} \sum_{j=1}^{d^{[l]}} \frac{\partial \text{loss}}{\partial z_j^{[l]}} \times \frac{\partial z_j^{[l]}}{\partial a_i^{[l-1]}} \\ &= f'(z_i^{[l-1]}) \sum_{j=1}^{d^{[l]}} \frac{\partial \text{loss}}{\partial z_j^{[l]}} \times w_{ji}^{[l]} \end{aligned}$$



Backpropagation

$$\frac{\partial \text{loss}}{\partial w_{ji}^{[l]}} = \boxed{\frac{\partial \text{loss}}{\partial z_j^{[l]}}} \times \frac{\partial z_j^{[l]}}{\partial w_{ji}^{[l]}}$$

$$= \boxed{\delta_j^{[l]}} \times a_i^{[l-1]}$$




$\delta_j^{[l]} = \frac{\partial \text{loss}}{\partial z_j^{[l]}}$ is the **sensitivity** of the loss to $z_j^{[l]}$

Sensitivity vectors can be obtained by running a backward process in the network architecture (hence the name backpropagation.)

We will compute $\delta^{[l-1]}$ from $\delta^{[l]}$:

$$\delta_i^{[l-1]} = f'(z_i^{[l-1]}) \sum_{j=1}^n \delta_j^{[l]} \times w_{ji}^{[l]}$$



Backpropagation Algorithm (Gradient Descent)

Initialize all weights to small random numbers.


While not satisfied

For each training example do:

1. Feed forward the training example to the network and compute the outputs of all units in forward step (z and a) and the loss

2. For each unit find its δ in the backward step

3. Update each network weight $w_{ij}^{[l]}$ as $w_{ij}^{[l]} \leftarrow w_{ij}^{[l]} - \eta \frac{\partial \text{loss}}{\partial w_{ij}^{[l]}}$ where $\frac{\partial \text{loss}}{\partial w_{ij}^{[l]}} = \delta_j^{[l]} \times a_i^{[l-1]}$



Next Session: Convolutional Neural Networks (CNNs)

